

EN3160: Image Processing and Computer Vision


Assignment 01: Intensity Transformations and Neighborhood Filtering

- **Name:** Ayanaja N.B.G.M
- **Index Number:** 230065H
- **GitHub Repository:** https://github.com/ManiiAya/EN3160_Assignment_01

Question 1: Piecewise Linear Intensity Transformations

```
In [106]: import cv2 as cv
import numpy as np
import matplotlib.pyplot as plt
from scipy.signal import convolve2d
```

```
In [107]: x = [0,50,50,100,150,150,255]
y = [0,50,100,170,255,150,255]
f = cv.imread("images/emma.jpg", cv.IMREAD_GRAYSCALE)
X = np.arange(256, dtype='uint8')
F = np.interp(X,x,y).astype(np.uint8)
fig, axes = plt.subplots(1,3, figsize=(24,8))
axes[0].imshow(f, cmap='gray', vmin=0, vmax=255)
axes[1].imshow(F[f], cmap='gray', vmin=0, vmax=255)
axes[2].plot(X,F)
plt.show()
```






```
In [108]: def intensity_transform(im, breakpoints):
x, y = zip(*breakpoints)
X = np.arange(256, dtype='uint8')
F = np.interp(X, x, y).astype(np.uint8)
return F[im], X, F
```

```
In [109]: f = cv.imread("images/emma.jpg", cv.IMREAD_GRAYSCALE)
breakpoints = [(0, 0), (64, 32), (128, 128), (192, 224), (255, 255)]
transformed_img, X, F = intensity_transform(f, breakpoints)
fig, axes = plt.subplots(1, 3, figsize=(15, 5))
axes[0].imshow(f, cmap='gray', vmin=0, vmax=255)
axes[0].set_title("Original Image");axes[0].xaxis.set_ticks_position('top')
axes[1].imshow(transformed_img, cmap='gray', vmin=0, vmax=255)
axes[1].set_title("Transformed Image");axes[1].xaxis.set_ticks_position('top') # Hides the axes
axes[2].plot(X, F, color='blue', label='Transformation Curve')
```

```

x_coords, y_coords = zip(*breakpoints)
axes[2].plot(x_coords, y_coords, 'ro', label='Breakpoints')
axes[2].set_title("Intensity Transformation")
axes[2].set_xlabel("Input Intensity")
axes[2].set_ylabel("Output Intensity")
axes[2].grid(True)
axes[2].legend()
plt.tight_layout()
plt.show()

```




Question 2: Feature Accentuation in Medical Imagery

```

In [110]: brain_img = cv.imread("images/brain.png", cv.IMREAD_GRAYSCALE)
hist_values, bin_edges, _ = plt.hist(brain_img.ravel(), bins=256, range=[0,256])
plt.show()
#printed the histogram values for better understanding of the distribution of pixel intensities
# for intensity in range(140,256):
#     print(f"Intensity {intensity}: {100*hist_values[intensity]/np.max(hist_values)}")

```



```

In [134]: f = cv.imread("images/brain.png", cv.IMREAD_GRAYSCALE)
breakpoints = [(0, 0), (153, 0), (154, 100), (168, 255), (184, 100), (185, 0), (221, 0)]

```

```

transformed_img, X, F = intensity_transform(f, breakpoints)
fig, axes = plt.subplots(1, 3, figsize=(15, 5))


axes[0].imshow(f, cmap='gray', vmin=0, vmax=255)
axes[0].set_title("Original Image")
axes[0].xaxis.set_ticks_position('top')

axes[1].imshow(transformed_img, cmap='gray', vmin=0, vmax=255)
axes[1].set_title("Accentuate White Matter")
axes[1].xaxis.set_ticks_position('top')

axes[2].plot(X, F, color='blue', label='Transformation Curve')
x_coors, y_coors = zip(*breakpoints)
axes[2].plot(x_coors, y_coors, 'ro', label='Breakpoints')
axes[2].set_title("Intensity Transformation")
axes[2].set_xlabel("Input Intensity")
axes[2].set_ylabel("Output Intensity")
axes[2].grid(True)
axes[2].legend()

plt.tight_layout()
plt.show()

```



```

In [135]: f = cv.imread("images/brain.png", cv.IMREAD_GRAYSCALE)
breakpoints = [(0, 0), (185, 0), (186, 100), (194, 255), (221, 100), (222, 0), (255, 0)]

transformed_img, X, F = intensity_transform(f, breakpoints)
fig, axes = plt.subplots(1, 3, figsize=(15, 5))


axes[0].imshow(f, cmap='gray', vmin=0, vmax=255)
axes[0].set_title("Original Image")
axes[0].xaxis.set_ticks_position('top')

axes[1].imshow(transformed_img, cmap='gray', vmin=0, vmax=255)
axes[1].set_title("Gray Matter")
axes[1].xaxis.set_ticks_position('top')

axes[2].plot(X, F, color='blue', label='Transformation Curve')
x_coors, y_coors = zip(*breakpoints)
axes[2].plot(x_coors, y_coors, 'ro', label='Breakpoints')
axes[2].set_title("Intensity Transformation")
axes[2].set_xlabel("Input Intensity")
axes[2].set_ylabel("Output Intensity")
axes[2].grid(True)
axes[2].legend()

```

```
plt.tight_layout()
plt.show()
```



Interpretation of Question 2: Tissue Accentuation

To analyze specific tissues in the brain proton density image, I first plotted the histogram to identify the distinct intensity clusters.

- **White Matter:** On the histogram, I identified this as the first major peak. I designed a triangular breakpoint filter ranging from 153 to 185, peaking at an intensity of **168**.
- **Gray Matter:** identified this as the second peak on the histogram and designed a triangular breakpoint filter ranging from 185 to 222, peaking at an intensity of **194**.


So by applying these custom breakpoint coordinates, the target matter is driven to 255 (white) while all other background and tissue values are suppressed to 0 (black).

Question 3: Gamma Correction in Lab Color Space

```
In [113]: img_original = cv.imread(r"D:\Semester 5\CV\EN3160_Assignment_01\images\white.jpg", cv.IMREAD_COLOR)
f = cv.cvtColor(img_original, cv.COLOR_BGR2LAB)
L,a,b = cv.split(f)
gamma = 0.5
Lc = (((L/255.0)**(gamma))*255).astype(np.uint8)
img_lab = cv.merge([Lc,a,b])
img = cv.cvtColor(img_lab, cv.COLOR_LAB2RGB)
fig, axes = plt.subplots(2, 2, figsize=(12, 9))
axes[0, 0].imshow(cv.cvtColor(img_original, cv.COLOR_BGR2RGB))
axes[0, 0].set_title("Original Image")
axes[0, 0].axis('off')
axes[0, 1].imshow(img)
axes[0, 1].set_title(f"Gamma Corrected Image (Gamma = {gamma})")
axes[0, 1].axis('off')
axes[1, 0].hist(L.ravel(), bins=256, range=[0, 256], color='gray')
axes[1, 0].set_title("Histogram: Original L Plane")
axes[1, 0].set_xlabel("Pixel Intensity")
axes[1, 0].set_ylabel("Frequency")
axes[1, 1].hist(Lc.ravel(), bins=256, range=[0, 256], color='gray')
axes[1, 1].set_title(f"Histogram: Corrected L Plane (Gamma = {gamma})")
axes[1, 1].set_xlabel("Pixel Intensity")
axes[1, 1].set_ylabel("Frequency")
```



```
plt.tight_layout()
plt.show()
```



Question 4: Vibrance Enhancement via Saturation Transformation

```
In [114]: f = cv.imread(r"D:\Semester 5\CV\EN3160_Assignment_01\images\spider.png", cv.IMREAD_COLOR)
img = cv.cvtColor(f, cv.COLOR_BGR2HSV)
h,s,v = cv.split(img)
sigma = 70
a = 0.45
i = np.arange(256)
F = i + a * (128 * np.exp(-((i-128)**2) / (2*sigma**2)))
F = np.clip(F,0,255).astype(np.uint8)
s_new = F[s]
img_new = cv.merge([h,s_new,v])
img_rgb = cv.cvtColor(img_new, cv.COLOR_HSV2RGB)
fig, axes = plt.subplots(1,3, figsize=(19,4))
axes[0].plot(i, F);axes[0].set_title("Vibrance Transformation Curve")
axes[0].set_xlabel("Input Saturation");axes[0].set_ylabel("Output Saturation")
axes[0].grid(True) # Adds a grid for better readability
axes[1].imshow(cv.cvtColor(f, cv.COLOR_BGR2RGB)); axes[1].xaxis.set_ticks_position('top') ;
axes[1].set_title("Original Image")
axes[2].imshow(img_rgb); axes[2].xaxis.set_ticks_position('top')
axes[2].set_title(f"Transformed Image - a: {a}")
plt.tight_layout()
plt.show()
```



To achieve a visually pleasing output that enhances the mid-tone colors of the image without oversaturating the bright reds, I adjusted the parameter to $a = 0.4$

Question 5: Global Histogram Equalization

```
In [115]: def equalize(f):
            hist, _ = np.histogram(f.flatten(), 256, [0, 256], True)
            cum_sum = np.cumsum(hist) * (255)
            F = cum_sum.astype(np.uint8)
            img = F[f]
            return img

f = cv.imread(r"D:\Semester 5\CV\EN3160_Assignment_01\images\shells.tif", cv.IMREAD_GRAYSCALE)
img_eq = equalize(f)
fig, axes = plt.subplots(1, 4, figsize=(20, 5))


# 1. Original Image
axes[0].imshow(f, cmap='gray', vmin=0, vmax=255)
axes[0].set_title("Original Image")
axes[0].axis("off")

# 2. Original Histogram
axes[1].hist(f.flatten(), bins=256, range=[0, 256], color='gray')
axes[1].set_title("Original Histogram")
axes[1].set_xlabel("Pixel Intensity")
axes[1].set_ylabel("Frequency")

# 3. Equalized Image
axes[2].imshow(img_eq, cmap='gray', vmin=0, vmax=255)
axes[2].set_title("Equalized Image")
axes[2].axis("off")

# 4. Equalized Histogram
axes[3].hist(img_eq.flatten(), bins=256, range=[0, 256], color='gray')
axes[3].set_title("Equalized Histogram")
axes[3].set_xlabel("Pixel Intensity")

plt.tight_layout()
plt.show()
```



Question 6: Foreground-Aware Histogram Equalization

```
In [116]: f = cv.imread(r"D:\Semester 5\CV\EN3160_Assignment_01\images\jeniffer.jpg", cv.IMREAD_COLOR)
h, s, v = cv.split(cv.cvtColor(f, cv.COLOR_BGR2HSV))
fig, axes = plt.subplots(1,3, figsize=(15,6))
axes[0].imshow(h, cmap='gray', vmin=0, vmax=255); axes[0].xaxis.set_ticks_position('top'); axes
axes[1].imshow(s, cmap='gray', vmin=0, vmax=255); axes[1].axis(False); axes[1].set_title("Satura
axes[2].imshow(v, cmap='gray', vmin=0, vmax=255); axes[2].axis(False); axes[2].set_title("Value"
plt.tight_layout()
plt.show()
```



By observing above three graphs, it can be concluded that the Saturation plane is optimal to mask the foreground.

```
In [117]: _, mask = cv.threshold(s, 11, 255, cv.THRESH_BINARY)
foreg_pixels = v[mask == 255]
hist, _ = np.histogram(foreg_pixels, 256, [0, 256], True)
cum_sum = np.cumsum(hist)*(255)
F = cum_sum.astype(np.uint8)
v_eq = v.copy()
v_eq[mask==255] = F[v[mask==255]]
img_rgb_eq = cv.cvtColor(cv.merge([h,s,v_eq]), cv.COLOR_HSV2RGB)
print(F.shape)
fig, axes = plt.subplots(1, 5, figsize=(22, 4))
axes[0].hist(v.flatten(), bins=256, color='gray')
axes[0].set_title("Original V Histogram")
axes[1].hist(foreg_pixels, bins=256, color='blue')
axes[1].set_title("Foreground Pixels Only")
axes[2].imshow(cv.cvtColor(f, cv.COLOR_BGR2RGB))
axes[2].set_title("Original Image")
axes[2].axis("off")
axes[3].imshow(img_rgb_eq)
axes[3].set_title("Foreground Equalized Image")
axes[3].axis("off")
axes[4].hist(v_eq.flatten(), bins=256, color='gray')
plt.tight_layout()
plt.show()
```

(256,)





Question 7: Edge Detection via Sobel Operator and Kernel Separability

```
In [118]: f = cv.imread(r"D:\Semester 5\CV\EN3160_Assignment_01\images\ein.png", cv.IMREAD_REDUCED_GRAYSCALE)
print(f.shape)
sobel_kernel = np.array([[ 1, 0, -1],
                          [ 2, 0, -2],
                          [ 1, 0, -1]])
img_filtered = cv.filter2D(f, cv.CV_64F, sobel_kernel)
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
axes[0].imshow(f, cmap='gray', vmin=0, vmax=255)
axes[0].set_title("Original Image")
axes[0].axis("off")

axes[1].imshow(img_filtered, cmap='gray')
axes[1].set_title("Sobel Filtered (Vertical Edges)")
axes[1].axis("off")

plt.tight_layout()
plt.show()
```

(190, 182)



```
In [119]: def sobel(img, kernel):
rows, cols = img.shape
output = np.zeros((rows, cols), dtype=np.float64)
pad_img = np.pad(img, pad_width=1, mode='constant', constant_values=0)
for i in range(rows):
    for j in range(cols):
        region = pad_img[i:i+3, j:j+3]
        output[i, j] = np.sum(region * kernel)
return output
```




```
f = cv.imread(r"D:\Semester 5\CV\EN3160_Assignment_01\images\ein.png", cv.IMREAD_GRAYSCALE)
sobel_kernel = np.array([[ 1, 0, -1],
                          [2, 0, -2],
                          [ 1, 0, -1]])
sobel_result = sobel(f, sobel_kernel)
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
axes[0].imshow(f, cmap='gray', vmin=0, vmax=255)
axes[0].set_title("Original Image")
axes[0].axis("off")


axes[1].imshow(sobel_result, cmap='gray')
axes[1].set_title("Sobel Filtered (Vertical Edges)")
axes[1].axis("off")

plt.tight_layout()
plt.show()
```


Original Image



Sobel Filtered (Vertical Edges)



```
In [120]: v_kernel = np.array([[1],
                               [2],
                               [1]])
h_kernel = np.array([[1, 0, -1]])
conv1 = cv.filter2D(f, cv.CV_64F, h_kernel)
img_filtered = cv.filter2D(conv1, cv.CV_64F, v_kernel)
plt.imshow(img_filtered, cmap='gray',)
plt.title("Separable 1D Filters")
plt.show()
```



Question 8: Image Scaling and Interpolation Evaluation

```
In [121]: def zoom_image(img, scale_factor, method='nearest'):
    h, w = img.shape[:2]
    new_h = int(h * scale_factor)
    new_w = int(w * scale_factor)
    r_new = np.arange(new_h)
    c_new = np.arange(new_w)
    r_old = r_new / scale_factor
    c_old = c_new / scale_factor

    if method == 'nearest':
        r_idx = np.clip(np.round(r_old).astype(int), 0, h - 1)
        c_idx = np.clip(np.round(c_old).astype(int), 0, w - 1)
        return img[r_idx[:, None], c_idx]

    elif method == 'bilinear':
        img_f = img.astype(np.float64)
        r0 = np.clip(np.floor(r_old).astype(int), 0, h - 1)
        r1 = np.clip(r0 + 1, 0, h - 1)
        c0 = np.clip(np.floor(c_old).astype(int), 0, w - 1)
        c1 = np.clip(c0 + 1, 0, w - 1)
        dr = (r_old - r0).reshape(-1, 1)
        dc = (c_old - c0).reshape(1, -1)
        if len(img.shape) == 3:
            dr = dr.reshape(-1, 1, 1)
            dc = dc.reshape(1, -1, 1)
        Ia = img_f[r0[:, None], c0] # Top-Left
        Ib = img_f[r0[:, None], c1] # Top-Right
        Ic = img_f[r1[:, None], c0] # Bottom-Left
        Id = img_f[r1[:, None], c1] # Bottom-Right
        wa = (1 - dr) * (1 - dc)
        wb = (1 - dr) * dc
```

```

        wc = dr * (1 - dc)
        wd = dr * dc
        zoomed_img = Ia * wa + Ib * wb + Ic * wc + Id * wd

    return np.clip(zoomed_img, 0, 255).astype(np.uint8)

else:
    raise ValueError("Method must be 'nearest' or 'bilinear'")

```

```

In [122]: image_pairs = [
    (r"D:\Semester 5\CV\EN3160_Assignment_01\images\im01small.png",
     r"D:\Semester 5\CV\EN3160_Assignment_01\images\im01.png"),

    (r"D:\Semester 5\CV\EN3160_Assignment_01\images\im02small.png",
     r"D:\Semester 5\CV\EN3160_Assignment_01\images\im02.png")
]

fig, axes = plt.subplots(2, 3, figsize=(18, 10))

scale_factor = 4

for i, (path_small, path_large) in enumerate(image_pairs):
    img_small = cv.cvtColor(cv.imread(path_small, cv.IMREAD_COLOR), cv.COLOR_BGR2RGB)
    img_large = cv.cvtColor(cv.imread(path_large, cv.IMREAD_COLOR), cv.COLOR_BGR2RGB)

    img_nn = zoom_image(img_small, scale_factor, method='nearest')
    img_bl = zoom_image(img_small, scale_factor, method='bilinear')

    ssd_nn = np.mean(np.square(img_large.astype(np.float64) - img_nn.astype(np.float64)))
    ssd_bl = np.mean(np.square(img_large.astype(np.float64) - img_bl.astype(np.float64)))

    axes[i, 0].imshow(img_large)
    axes[i, 0].set_title(f"Pair {i+1}: Original High-Res")
    axes[i, 0].axis("off")

    axes[i, 1].imshow(img_nn)
    axes[i, 1].set_title(f"Nearest-Neighbor (MSD: {ssd_nn:.4f})")
    axes[i, 1].axis("off")

    axes[i, 2].imshow(img_bl)
    axes[i, 2].set_title(f"Bilinear (MSD: {ssd_bl:.4f})")
    axes[i, 2].axis("off")


plt.tight_layout()
plt.show()

```



Question 9: GrabCut Segmentation and Depth-of-Field Enhancement

```
In [123]: img_flower = cv.imread(r"D:\Semester 5\CV\EN3160_Assignment_01\images\daisy.jpg")
img_rgb = cv.cvtColor(img_flower, cv.COLOR_BGR2RGB)
mask = np.zeros(img_flower.shape[:2], np.uint8)
bgdModel = np.zeros((1, 65), np.float64)
fgdModel = np.zeros((1, 65), np.float64)
rect = (20, 130, 535, 470)
cv.grabCut(img_flower, mask, rect, bgdModel, fgdModel, 5, cv.GC_INIT_WITH_RECT)
mask_binary = np.where((mask == 2) | (mask == 0), 0, 1).astype('uint8')
foreground = cv.bitwise_and(img_rgb, img_rgb, mask=mask_binary)
mask_inv = cv.bitwise_not(mask_binary * 255) # Invert mask for background
background = cv.bitwise_and(img_rgb, img_rgb, mask=mask_inv)
fig, axes = plt.subplots(1, 3, figsize=(15, 5))
axes[0].imshow(mask_binary, cmap='gray'); axes[0].set_title("Segmentation Mask"); axes[0].axis('off')
axes[1].imshow(foreground); axes[1].set_title("Foreground"); axes[1].axis('off')
axes[2].imshow(background); axes[2].set_title("Background"); axes[2].axis('off')
plt.tight_layout()
plt.show()
```



```
In [125]: blurred_img = cv.GaussianBlur(img_rgb, (51, 51), 0)
mask_3d = mask_binary[:, :, np.newaxis]
```



```

enhanced_img = np.where(mask_3d == 1, img_rgb, blurred_img)
fig, axes = plt.subplots(1, 2, figsize=(12, 6))
axes[0].imshow(img_rgb); axes[0].set_title("Original Image"); axes[0].axis('off')
axes[1].imshow(enhanced_img); axes[1].set_title("Enhanced (Portrait Mode)"); axes[1].axis('off')
plt.tight_layout()
plt.show()

```

Original Image



Enhanced (Portrait Mode)



9(c) Discussion

The dark region beyond the edge of the flower in the enhanced image is an issue caused by the spatial filtering process. It's called the "color bleeding" or "halo effect".

When the heavy Gaussian blur is applied to the original image to create the artificial depth-of-field, the filter calculates a weighted average of neighboring pixels. Along the boundary of the flower, the dark pixels from the flower's edges and internal shadows are averaged together with the adjacent lighter background pixels. This spatial mixing artificially darkens the background pixels immediately surrounding the flower.

When the sharp, segmented foreground mask is superimposed back onto this blurred image, the darkened background pixels just outside the segmentation boundary remain visible, resulting in the dark halo effect.

Question 10: Edge-Preserving Smoothing via Bilateral Filtering


```

In [126]: img_lake = cv.imread(r"D:\Semester 5\CV\EN3160_Assignment_01\images\minaret.jpg", cv.IMREAD_COLOR)
img_rgb = cv.cvtColor(img_lake, cv.COLOR_BGR2RGB)
img_gaussian = cv.GaussianBlur(img_rgb, (9, 9), sigmaX=1)
img_bilateral_cv = cv.bilateralFilter(img_rgb, d=9, sigmaColor=75, sigmaSpace=15)

```



```
fig, axes = plt.subplots(1, 3, figsize=(18,6))
axes[0].imshow(img_rgb); axes[0].set_title("Original Image"); axes[0].axis('off')
axes[1].imshow(img_gaussian); axes[1].set_title("Gaussian Blur"); axes[1].axis('off')
axes[2].imshow(img_bilateral_cv); axes[2].set_title("OpenCV Bilateral"); axes[2].axis('off')
plt.tight_layout()
plt.show()
```



```
In [127]: def custom_bilateral_filter_channel(img, d, sigma_color, sigma_space):
    """Applies bilateral filter to a single 2D channel."""
    img = img.astype(np.float64)
    rows, cols = img.shape
    output = np.zeros((rows, cols), dtype=np.float64)
    r = d // 2
    x, y = np.meshgrid(np.arange(-r, r+1), np.arange(-r, r+1))
    spatial_weights = np.exp(-(x**2 + y**2) / (2 * sigma_space**2))
    padded_img = np.pad(img, pad_width=r, mode='reflect')
    for i in range(rows):
        for j in range(cols):
            patch = padded_img[i:i+d, j:j+d]
            center_val = padded_img[i+r, j+r]
            range_weights = np.exp(-((patch - center_val)**2) / (2 * sigma_color**2))
            total_weights = spatial_weights * range_weights
            output[i, j] = (np.sum(total_weights * patch) / np.sum(total_weights)).astype(np.float64)

    return output

def custom_bilateral_color(img_rgb, d, sigma_color, sigma_space):
    """Splits color image, filters channels, and merges."""
    r, g, b = cv.split(img_rgb)
    r_filtered = custom_bilateral_filter_channel(r, d, sigma_color, sigma_space)
    g_filtered = custom_bilateral_filter_channel(g, d, sigma_color, sigma_space)
    b_filtered = custom_bilateral_filter_channel(b, d, sigma_color, sigma_space)

    merged = cv.merge([r_filtered, g_filtered, b_filtered])
    return np.clip(merged, 0, 255).astype(np.uint8)
```

```
In [128]: img_bilateral_custom = custom_bilateral_color(img_rgb, d=9, sigma_color=75, sigma_space=15)
mse = np.mean((img_bilateral_cv.astype("float") - img_bilateral_custom.astype("float")) ** 2)
print(f"Mean Squared Error (OpenCV vs Custom): {mse:.4f}")

fig, axes = plt.subplots(1, 4, figsize=(24, 6))
axes[0].imshow(img_rgb)
axes[0].set_title("Original Image", fontsize=14)
axes[0].axis('off')
axes[1].imshow(img_gaussian)
axes[1].set_title("Gaussian Blur", fontsize=14)
axes[1].axis('off')
```


```

axes[2].imshow(img_bilateral_custom)
axes[2].set_title("Custom Bilateral Filter", fontsize=14)
axes[2].axis('off')
axes[3].imshow(img_bilateral_cv)
axes[3].set_title(f"OpenCV Bilateral (MSE vs Custom: {mse:.2f})", fontsize=14)
axes[3].axis('off')

plt.tight_layout()
plt.show()

```

Mean Squared Error (OpenCV vs Custom): 118.5314



10(b) Discussion: Custom vs. OpenCV Bilateral Filter

Custom bilateral filter has smoothed while preserving the edges to some extent. It has occurred a painterly effect. The mean squared error is 118.5 .

This larger MSE created due to the way that the R,G,B color channels are handled. In the custom method, it processes these channels separately and calculat the range weights. This creates minor color bleeding variations near edges. But OpenCV's bilateral filter calculutes the internsity difference using L_2 norm of the 3 channel color vectors and apply unified weuight which perfectly preserves the original hue.